<table>
<tr><td rowspan="2">cases</td><th>doc_1</th><th>doc_2</th><th>decision</th><th>id</th></tr>
<tr><td>
<table>
<tr><th>authors</th><td>Duncan, A.
NÃ¼sken, N.
Szpruch, L.</td></tr>
<tr><th>title</th><td>On the geometry of Stein variational gradient descent</td></tr>
<tr><th>publication_date</th><td>2019-12-01 00:00:00</td></tr>
<tr><th>source</th><td>SupportedSources.CORE</td></tr>
<tr><th>journal</th><td></td></tr>
<tr><th>volume</th><td></td></tr>
<tr><th>doi</th><td>None</td></tr>
<tr><th>urls</th><td>https://core.ac.uk/download/267952006.pdf</td></tr>
<tr><th>id</th><td>id-5044354038293896764</td></tr>
<tr><th>abstract</th><td>Bayesian inference problems require sampling or approximating high-dimensional probability distributions. The focus of this paper is on the recently introduced Stein variational gradient descent methodology, a class of algorithms that rely on iterated steepest descent steps with respect to a reproducing kernel Hilbert space norm. This construction leads to interacting particle systems, the mean-field limit of which is a gradient flow on the space of probability distributions equipped with a certain geometrical structure. We leverage this viewpoint to shed some light on the convergence properties of the algorithm, in particular addressing the problem of choosing a suitable positive definite kernel function. Our analysis leads us to considering certain nondifferentiable kernels with adjusted tails. We demonstrate significant performs gains of these in various numerical experiments</td></tr>
<tr><th>versions</th><td></td></tr>
</table>
</td><td>
<table>
<tr><th>authors</th><td>Duncan, A.
Nuesken, N.
Szpruch, L.</td></tr>
<tr><th>title</th><td>On the geometry of Stein variational gradient descent</td></tr>
<tr><th>publication_date</th><td>2019-12-02 00:00:00</td></tr>
<tr><th>source</th><td>SupportedSources.CORE</td></tr>
<tr><th>journal</th><td></td></tr>
<tr><th>volume</th><td></td></tr>
<tr><th>doi</th><td>None</td></tr>
<tr><th>urls</th><td>http://arxiv.org/abs/1912.00894</td></tr>
<tr><th>id</th><td>id567896283371958501</td></tr>
<tr><th>abstract</th><td>Bayesian inference problems require sampling or approximating high-dimensional probability distributions. The focus of this paper is on the recently introduced Stein variational gradient descent methodology, a class of algorithms that rely on iterated steepest descent steps with respect to a reproducing kernel Hilbert space norm. This construction leads to interacting particle systems, the mean-field limit of which is a gradient flow on the space of probability distributions equipped with a certain geometrical structure. We leverage this viewpoint to shed some light on the convergence properties of the algorithm, in particular addressing the problem of choosing a suitable positive definite kernel function. Our analysis leads us to considering certain nondifferentiable kernels with adjusted tails. We demonstrate significant performs gains of these in various numerical experiments.Comment: 39 pages, 4 figure</td></tr>
<tr><th>versions</th><td></td></tr>
</table>
</td><td>DUPLICATES</td><td>827</td></tr>
</table>