

cases	doc_1		doc_2		decision	id
	authors	<ul style="list-style-type: none">Chen, WenqiZhang, RuihanZhang, XiajieZhang, Yuhan	authors	<ul style="list-style-type: none">Chen, WenqiZhang, RuihanZhang, XiajieZhang, Yuhan	DUPLICATES	5
	title	Representing affect information in word embeddings	title	Representing Affect Information in Word Embeddings		
	publication_date	2023-01-27 00:00:00	publication_date	2022-09-21 00:00:00		
	source	SupportedSources.CORE	source	SupportedSources.CORE		
	journal		journal			
	volume		volume			
	doi	10.3765/elm.2.5391	doi	None		
	urls	<ul style="list-style-type: none">https://core.ac.uk/download/553281587.pdf	urls	<ul style="list-style-type: none">http://arxiv.org/abs/2209.10583		
	id	id8957577934454200204	id	id3978004612439189321		
	abstract	A growing body of research in natural language processing (NLP) and natural language understanding (NLU) is investigating human-like knowledge learned or encoded in the word embeddings from large language models. This is a step towards understanding what knowledge language models capture that resembles human understanding of language and communication. Here, we investigated whether and how the affect meaning of a word (i.e., valence, arousal, dominance) is encoded in word embeddings pre-trained in large neural networks. We used the human-labeled dataset (Mohammad 2018) as the ground truth and performed various correlational and classification tests on four types of word embeddings. The embeddings varied in being static or contextualized, and how much affect specific information was prioritized during the pre-training and fine-tuning phase. Our analyses show that word embedding from the vanilla BERT model (Devlin et al. 2019) did not saliently encode the affect information of English words. Only when the BERT model was fine-tuned on emotion related tasks or contained extra contextualized information from emotion-rich contexts could the corresponding embedding encode more relevant affect information	abstract	A growing body of research in natural language processing (NLP) and natural language understanding (NLU) is investigating human-like knowledge learned or encoded in the word embeddings from large language models. This is a step towards understanding what knowledge language models capture that resembles human understanding of language and communication. Here, we investigated whether and how the affect meaning of a word (i.e., valence, arousal, dominance) is encoded in word embeddings pre-trained in large neural networks. We used the human-labeled dataset as the ground truth and performed various correlational and classification tests on four types of word embeddings. The embeddings varied in being static or contextualized, and how much affect specific information was prioritized during the pre-training and fine-tuning phase. Our analyses show that word embedding from the vanilla BERT model did not saliently encode the affect information of English words. Only when the BERT model was fine-tuned on emotion-related tasks or contained extra contextualized information from emotion-rich contexts could the corresponding embedding encode more relevant affect information		
	versions		versions			