cases	doc_1		doc_2		decision	id
	authors	• Andrey Kutuzov • Mario Giulianelli	**authors**	• Andrey Kutuzov • Mario Giulianelli	DUPLICATES	48
	title	UiO-UvA at SemEval-2020 Task 1: Contextualised Embeddings for Lexical Semantic Change Detection	**title**	UiO-UvA at SemEval-2020 Task 1: Contextualised Embeddings for Lexical Semantic Change Detection		
	publication_date	2020-04-30 00:00:00	**publication_date**	2020-04-30 18:43:57+00:00		
	source	SupportedSources.SEMANTIC_SCHOLAR	**source**	SupportedSources.ARXIV		
	journal		**journal**	None		
	volume		**volume**			
	doi	10.18653/v1/2020.semeval-1.14	**doi**			
	urls	• https://www.semanticscholar.org/paper/77c1fd8dd9de152168f61420ba00310ce6492450	**urls**	• http://arxiv.org/pdf/2005.00050v3 • http://arxiv.org/abs/2005.00050v3 • http://arxiv.org/pdf/2005.00050v3		
	id	id-341829475772431587	**id**	id-1264066212405507433		
	abstract	We apply contextualised word embeddings to lexical semantic change detection in the SemEval-2020 Shared Task 1. This paper focuses on Subtask 2, ranking words by the degree of their semantic drift over time. We analyse the performance of two contextualising architectures (BERT and ELMo) and three change detection algorithms. We find that the most effective algorithms rely on the cosine similarity between averaged token embeddings and the pairwise distances between token embeddings. They outperform strong baselines by a large margin (in the post-evaluation phase, we have the best Subtask 2 submission for SemEval-2020 Task 1), but interestingly, the choice of a particular algorithm depends on the distribution of gold scores in the test set.	**abstract**	We apply contextualised word embeddings to lexical semantic change detection in the SemEval-2020 Shared Task 1. This paper focuses on Subtask 2, ranking words by the degree of their semantic drift over time. We analyse the performance of two contextualising architectures (BERT and ELMo) and three change detection algorithms. We find that the most effective algorithms rely on the cosine similarity between averaged token embeddings and the pairwise distances between token embeddings. They outperform strong baselines by a large margin (in the post-evaluation phase, we have the best Subtask 2 submission for SemEval-2020 Task 1), but interestingly, the choice of a particular algorithm depends on the distribution of gold scores in the test set.		
	versions		**versions**			