

cases	doc_1		doc_2				decision	id
							DUPLICATES	209
	authors	<ul style="list-style-type: none">Andrei-Marius AvramD. CatrinaDumitru-Clementin CercelMihai DascaluTraian RebedeaV. PaisDan Tufis	authors	<ul style="list-style-type: none">Andrei-Marius AvramDarius CatrinaDumitru-Clementin CercelMihai DascualuTraian RebedeaVasile PuaicsDan Tufics				
	title	Distilling the Knowledge of Romanian BERTs Using Multiple Teachers	title	Distilling the Knowledge of Romanian BERTs Using Multiple Teachers				
	publication_date	2021-12-23 00:00:00	publication_date	2021-12-23 00:00:00				
	source	SupportedSources.OPENALEX	source	SupportedSources.SEMANTIC_SCHOLAR				
	journal	arXiv (Cornell University)	journal					
	volume		volume					
	doi	10.48550/arxiv.2112.12650	doi					
	urls	<ul style="list-style-type: none">https://openalex.org/W4226067207https://doi.org/10.48550/arxiv.2112.12650http://arxiv.org/pdf/2112.12650	urls	<ul style="list-style-type: none">https://www.semanticscholar.org/paper/c2848530890acd7549fe7cfec7d4fc5a80c2feb7				
	id	id-3114794533438077014	id	id4075326683460524351				
	abstract		abstract	Running large-scale pre-trained language models in computationally constrained environments remains a challenging problem yet to be addressed, while transfer learning from these models has become prevalent in Natural Language Processing tasks. Several solutions, including knowledge distillation, network quantization, or network pruning have been previously proposed; however, these approaches focus mostly on the English language, thus widening the gap when considering low-resource languages. In this work, we introduce three light and fast versions of distilled BERT models for the Romanian language: Distil-BERT-base-ro, Distil-RoBERT-base, and DistilMulti-BERT-base-ro. The first two models resulted from the individual distillation of knowledge from two base versions of Romanian BERTs available in literature, while the last one was obtained by distilling their ensemble. To our knowledge, this is the first attempt to create publicly available Romanian distilled BERT models, which were thoroughly evaluated on five tasks: part-of-speech tagging, named entity recognition, sentiment analysis, semantic textual similarity, and dialect identification. Our experimental results argue that the three distilled models offer performance comparable to their teachers, while being twice as fast on a GPU and ~35% smaller. In addition, we further test the similarity between the predictions of our students versus their teachers by measuring their label and probability loyalty, together with regression loyalty - a new metric introduced in this work.				
	versions		versions					