

| | | | | | | |
|-------|------------------|--|-------|--|------------|-----|
| cases | doc_1 | | doc_2 | | decision | id |
| | authors | <ul style="list-style-type: none">Jui ShahYaman Kumar SinglaChangyou ChenRajiv Ratn Shah | | | DUPLICATES | 232 |
| | title | What all do audio transformer models hear? Probing Acoustic Representations for Language Delivery and its Structure | | | | |
| | publication_date | 2021-01-02 06:29:12+00:00 | | | | |
| | source | SupportedSources.ARXIV | | | | |
| | journal | None | | | | |
| | volume | | | | | |
| | doi | | | | | |
| | urls | <ul style="list-style-type: none">http://arxiv.org/pdf/2101.00387v2http://arxiv.org/abs/2101.00387v2http://arxiv.org/pdf/2101.00387v2 | | | | |
| | id | id-7345381725309213464 | | | | |
| | abstract | In recent times, BERT based transformer models have become an inseparable part of the 'tech stack' of text processing models. Similar progress is being observed in the speech domain with a multitude of models observing state-of-the-art results by using audio transformer models to encode speech. This begs the question of what are these audio transformer models learning. Moreover, although the standard methodology is to choose the last layer embedding for any downstream task, but is it the optimal choice? We try to answer these questions for the two recent audio transformer models, Mockingjay and wave2vec2.0. We compare them on a comprehensive set of language delivery and structure features including audio, fluency and pronunciation features. Additionally, we probe the audio models' understanding of textual surface, syntax, and semantic features and compare them to BERT. We do this over exhaustive settings for native, non-native, synthetic, read and spontaneous speech datasets | | | | |
| | versions | | | | | |
| | authors | <ul style="list-style-type: none">Jui ShahYaman Kumar SinglaChangyou ChenR. Shah | | | | |
| | title | What all do audio transformer models hear? Probing Acoustic Representations for Language Delivery and its Structure | | | | |
| | publication_date | 2021-01-02 00:00:00 | | | | |
| | source | SupportedSources.SEMANTIC_SCHOLAR | | | | |
| | journal | ArXiv | | | | |
| | volume | abs/2101.00387 | | | | |
| | doi | 10.20944/preprints202101.0081.v1 | | | | |
| | urls | <ul style="list-style-type: none">https://www.semanticscholar.org/paper/5e32055b61bf58c336af41ca8642c7138a753aa7 | | | | |
| | id | id3287545049555777753 | | | | |
| | abstract | In recent times, BERT based transformer models have become an inseparable part of the “tech stack” of text processing models. Similar progress is being observed in the speech domain with a multitude of models observing state-of-the-art results by using audio transformer models to encode speech. This begs the question of what are these audio transformer models learning. Moreover, although the standard methodology is to choose the last layer embeddings for any downstream task, but is it the optimal choice? We try to answer these questions for the two recent audio transformer models, Mockingjay and wave2vec2.0 . We compare them on a comprehensive set of language delivery and structure features including audio, fluency and pronunciation features. Additionally, we probe the audio models’s understanding of textual surface, syntax, and semantic features and compare them to BERT. We do this over exhaustive settings for native, nonnative, synthetic, read and spontaneous speech datasets. | | | | |
| | versions | | | | | |