

cases	doc_1		doc_2		decision	id
	authors	<ul style="list-style-type: none">Jianmo NiGustavo Hern'andez 'AbregoNoah ConstantJi MaKeith B. HallDaniel Matthew CerYinfei Yang	authors	<ul style="list-style-type: none">Jianmo NiGustavo HernÃ­ndez ÃbregoNoah ConstantJi MaKeith B. HallDaniel CerYinfei Yang	DUPLICATES	247
	title	Sentence-T5: Scalable Sentence Encoders from Pre-trained Text-to-Text Models	title	Sentence-T5: Scalable Sentence Encoders from Pre-trained Text-to-Text Models		
	publication_date	2021-08-19 00:00:00	publication_date	2021-12-14 00:00:00		
	source	SupportedSources.SEMANTIC_SCHOLAR	source	SupportedSources.INTERNET_ARCHIVE		
	journal	ArXiv	journal			
	volume	abs/2108.08877	volume			
	doi	10.18653/v1/2022.findings-acl.146	doi			
	urls	<ul style="list-style-type: none">https://www.semanticscholar.org/paper/dbe87b171bfb789e1d22a047aeeee69105e6fd02	urls	<ul style="list-style-type: none">https://web.archive.org/web/20211220012551/https://arxiv.org/pdf/2108.08877v3.pdf		
	id	id-8907498693760694982	id	id8550156269653709715		
	abstract	We provide the first exploration of sentence embeddings from text-to-text transformers (T5) including the effects of scaling up sentence encoders to 11B parameters. Sentence embeddings are broadly useful for language processing tasks. While T5 achieves impressive performance on language tasks, it is unclear how to produce sentence embeddings from encoder-decoder models. We investigate three methods to construct Sentence-T5 (ST5) models: two utilize only the T5 encoder and one using the full T5 encoder-decoder. We establish a new sentence representation transfer benchmark, SentGLUE, which extends the SentEval toolkit to nine tasks from the GLUE benchmark. Our encoder-only models outperform the previous best models on both SentEval and SentGLUE transfer tasks, including semantic textual similarity (STS). Scaling up ST5 from millions to billions of parameters shown to consistently improve performance. Finally, our encoder-decoder method achieves a new state-of-the-art on STS when using sentence embeddings.	abstract	We provide the first exploration of sentence embeddings from text-to-text transformers (T5). Sentence embeddings are broadly useful for language processing tasks. While T5 achieves impressive performance on language tasks cast as sequence-to-sequence mapping problems, it is unclear how to produce sentence embeddings from encoder-decoder models. We investigate three methods for extracting T5 sentence embeddings: two utilize only the T5 encoder and one uses the full T5 encoder-decoder model. To support our investigation, we establish a new sentence representation transfer benchmark, SentGLUE, which extends the SentEval toolkit to nine tasks from the GLUE benchmark. Our encoder-only models outperforms Sentence-BERT and SimCSE sentence embeddings on both SentEval and SentGLUE transfer tasks, including semantic textual similarity (STS). Scaling up T5 from millions to billions of parameters is found to produce consistent further improvements. Finally, our encoder-decoder method achieves a new state-of-the-art on STS when using sentence embeddings. Our models are released at https://tfhub.dev/google/collections/sentence-t5/1.		
	versions		versions			