

cases	doc_1		doc_2		decision	id
					DUPLICATES	249
	authors	<ul style="list-style-type: none">Bohan LiHao ZhouJunxian HeMingxuan WangYiming YangLei Li	authors	<ul style="list-style-type: none">Lei LiYiming YangMingxuan WangJunxian HeHao ZhouBohan Li		
	title	On the Sentence Embeddings from Pre-trained Language Models	title	On the Sentence Embeddings from Pre-trained Language Models		
	publication_date	2020-11-24 00:00:00	publication_date	2020-11-02 00:00:00		
	source	SupportedSources.OPENALEX	source	SupportedSources.PAPERS_WITH_CODE		
	journal	Empirical Methods in Natural Language Processing	journal			
	volume		volume			
	doi	10.18653/v1/2020.emnlp-main.733	doi			
	urls	<ul style="list-style-type: none">https://openalex.org/W3105816068https://doi.org/10.18653/v1/2020.emnlp-main.733https://www.aclweb.org/anthology/2020.emnlp-main.733.pdf	urls	<ul style="list-style-type: none">https://arxiv.org/pdf/2011.05864v1.pdfhttps://github.com/InsaneLife/dssmhttps://aclanthology.org/2020.emnlp-main.733.pdf		
	id	id-3200391722590569090	id	id6959606668887333883		
	abstract		abstract	Pre-trained contextual representations like BERT have achieved great success in natural language processing. However, the sentence embeddings from the pre-trained language models without fine-tuning have been found to poorly capture semantic meaning of sentences. In this paper, we argue that the semantic information in the BERT embeddings is not fully exploited. We first reveal the theoretical connection between the masked language model pre-training objective and the semantic similarity task theoretically, and then analyze the BERT sentence embeddings empirically. We find that BERT always induces a non-smooth anisotropic semantic space of sentences, which harms its performance of semantic similarity. To address this issue, we propose to transform the anisotropic sentence embedding distribution to a smooth and isotropic Gaussian distribution through normalizing flows that are learned with an unsupervised objective. Experimental results show that our proposed BERT-flow method obtains significant performance gains over the state-of-the-art sentence embeddings on a variety of semantic textual similarity tasks. The code is available at https://github.com/bohanli/BERT-flow.		
	versions		versions			