

cases	doc_1		doc_2				decision	id		
							DUPLICATES	261		
	authors	<ul style="list-style-type: none">Vladimir AraujoAndrés CarvalloCarolina AspillagaDenis Parra	authors	<ul style="list-style-type: none">Vladimir AraujoAndrés CarvalloC. AspillagaDenis Parra						
			title						On Adversarial Examples for Biomedical NLP Tasks	
	publication_date								2020-04-23 00:00:00	
	title	On Adversarial Examples for Biomedical NLP Tasks		source	SupportedSources.SEMANTIC_SCHOLAR					
	publication_date	2020-04-23 00:00:00		journal	ArXiv					
	source	SupportedSources.OPENALEX		volume	abs/2004.11157					
	journal	arXiv (Cornell University)		doi						
	volume			urls	<ul style="list-style-type: none">https://www.semanticscholar.org/paper/6f509be31111ed828486198ad3ca41acb5c1c257					
	doi	None		id	id4600450713927831273					
	urls	<ul style="list-style-type: none">https://openalex.org/W3017885892		abstract	The success of pre-trained word embeddings has motivated its use in tasks in the biomedical domain. The BERT language model has shown remarkable results on standard performance metrics in tasks such as Named Entity Recognition (NER) and Semantic Textual Similarity (STS), which has brought significant progress in the field of NLP. However, it is unclear whether these systems work seemingly well in critical domains, such as legal or medical. For that reason, in this work, we propose an adversarial evaluation scheme on two well-known datasets for medical NER and STS. We propose two types of attacks inspired by natural spelling errors and typos made by humans. We also propose another type of attack that uses synonyms of medical terms. Under these adversarial settings, the accuracy of the models drops significantly, and we quantify the extent of this performance loss. We also show that we can significantly improve the robustness of the models by training them with adversarial examples. We hope our work will motivate the use of adversarial examples to evaluate and develop models with increased robustness for medical tasks.					
	id	id-6493339733811671171		versions						
	abstract									
	versions									