cases

doc_1		doc_2		decision	id
authors	• Ashwin K Vijayakumar • Ramakrishna Vedantam • Devi Parikh	**authors**	• Ashwin K. Vijayakumar • Ramakrishna Vedantam • Devi Parikh	DUPLICATES	372
title	Sound-Word2Vec: Learning Word Representations Grounded in Sounds	**title**	Sound-Word2Vec: Learning Word Representations Grounded in Sounds		
publication_date	2017-03-06 04:30:12+00:00	**publication_date**	2017-03-06 00:00:00		
source	SupportedSources.ARXIV	**source**	SupportedSources.SEMANTIC_SCHOLAR		
journal	None	**journal**			
volume		**volume**			
doi		**doi**	10.18653/v1/D17-1096		
urls	• http://arxiv.org/pdf/1703.01720v4 • http://arxiv.org/abs/1703.01720v4 • http://arxiv.org/pdf/1703.01720v4	**urls**	• https://www.semanticscholar.org/paper/73deebd57b5abb7b26a0ff6cbd6b13f2fec45501		
id	id-2442767175628276178	**id**	id-5665529979687597632		
abstract	To be able to interact better with humans, it is crucial for machines to understand sound - a primary modality of human perception. Previous works have used sound to learn embeddings for improved generic textual similarity assessment. In this work, we treat sound as a first-class citizen, studying downstream textual tasks which require aural grounding. To this end, we propose sound-word2vec - a new embedding scheme that learns specialized word embeddings grounded in sounds. For example, we learn that two seemingly (semantically) unrelated concepts, like leaves and paper are similar due to the similar rustling sounds they make. Our embeddings prove useful in textual tasks requiring aural reasoning like text-based sound retrieval and discovering foley sound effects (used in movies). Moreover, our embedding space captures interesting dependencies between words and onomatopoeia and outperforms prior work on aurally-relevant word relatedness datasets such as AMEN and ASLex.	**abstract**	To be able to interact better with humans, it is crucial for machines to understand sound â€“ a primary modality of human perception. Previous works have used sound to learn embeddings for improved generic semantic similarity assessment. In this work, we treat sound as a first-class citizen, studying downstream 6textual tasks which require aural grounding. To this end, we propose sound-word2vec â€“ a new embedding scheme that learns specialized word embeddings grounded in sounds. For example, we learn that two seemingly (semantically) unrelated concepts, like leaves and paper are similar due to the similar rustling sounds they make. Our embeddings prove useful in textual tasks requiring aural reasoning like text-based sound retrieval and discovering Foley sound effects (used in movies). Moreover, our embedding space captures interesting dependencies between words and onomatopoeia and outperforms prior work on aurally-relevant word relatedness datasets such as AMEN and ASLex.		
versions		**versions**			