cases	doc_1		doc_2		decision	id
	authors	• Aditya Balu • Sergio Botelho • Biswajit Khara • Vinay V Rao • Soumik Sarkar • Chinmay Hegde • Adarsh Krishnamurthy • Santi Adavani • Baskar Ganapathysubramanian	**authors**	• Aditya Balu • Sergio Botelho • Biswajit Khara • Vinay Rao • Chinmay Hegde • Soumik Sarkar • Santi Adavani • Adarsh Krishnamurthy • Baskar Ganapathysubramanian	DUPLICATES	212
	title	Distributed multigrid neural solvers on megavoxel domains	**title**	Distributed Multigrid Neural Solvers on Megavoxel Domains		
	publication_date	2021-11-13 00:00:00	**publication_date**	2021-04-29 00:00:00		
	source	SupportedSources.OPENALEX	**source**	SupportedSources.INTERNET_ARCHIVE		
	journal	IEEE International Conference on High Performance Computing, Data, and Analytics	**journal**			
	volume		**volume**			
	doi	10.1145/3458817.3476218	**doi**			
	urls	• https://openalex.org/W3206964070 • https://doi.org/10.1145/3458817.3476218 • https://dl.acm.org/doi/pdf/10.1145/3458817.3476218	**urls**	• https://web.archive.org/web/20210502084252/https://arxiv.org/pdf/2104.14538v1.pdf		
	id	id3429805577296147812	**id**	id1440328196587944127		
	abstract		**abstract**	We consider the distributed training of large-scale neural networks that serve as PDE solvers producing full field outputs. We specifically consider neural solvers for the generalized 3D Poisson equation over megavoxel domains. A scalable framework is presented that integrates two distinct advances. First, we accelerate training a large model via a method analogous to the multigrid technique used in numerical linear algebra. Here, the network is trained using a hierarchy of increasing resolution inputs in sequence, analogous to the 'V', 'W', 'F', and 'Half-V' cycles used in multigrid approaches. In conjunction with the multi-grid approach, we implement a distributed deep learning framework which significantly reduces the time to solve. We show the scalability of this approach on both GPU (Azure VMs on Cloud) and CPU clusters (PSC Bridges2). This approach is deployed to train a generalized 3D Poisson solver that scales well to predict output full-field solutions up to the resolution of 512x512x512 for a high dimensional family of inputs.		
	versions		**versions**			